

# I/O Management

# Categories of I/O Devices

## ✓ Human readable

- used to communicate with the user
- video display terminals
- keyboard
- mouse
- printer

# Categories of I/O Devices

## ✓ Machine readable

- used to communicate with electronic equipment
- disk drives
- tape drives
- controllers
- actuators

# Categories of I/O Devices

## ✓ Communication

- used to communicate with remote devices
- digital line drivers
- modems

# Differences in I/O Devices

## ✓ Data Transfer Rate

## ✓ Application-specific

- disk used to store files must have file-management software
- disk used to store virtual memory pages depends on virtual memory hardware; I/O ops may be scheduled differently than for disks used for file storage
- terminal used by system administrator may have a higher priority

# Differences in I/O Devices

- ✓ Complexity of control
- ✓ Unit of transfer
  - data may be transferred as a stream of bytes for a terminal or in larger blocks for a disk
- ✓ Data representation
  - encoding schemes: character encoding, parity may be different
- ✓ Error conditions
  - devices respond to errors differently

# Techniques for Performing I/O

## ✓ Programmed I/O

- process is busy-waiting for the operation to complete

## ✓ Interrupt-driven I/O

- I/O command is issued
- processor continues executing instructions
- I/O module sends an interrupt when done

# Techniques for Performing I/O

## ✓ Direct Memory Access (DMA)

- DMA module controls exchange of data between main memory and the I/O device
- processor interrupted only after entire block has been transferred



# Evolution of the I/O Function

- ✓ Processor directly controls a peripheral device
- ✓ Controller or I/O module is added
  - processor uses programmed I/O without interrupts
  - processor does not need to handle details of external devices

# Evolution of the I/O Function

- ✓ Controller or I/O module with interrupts
  - processor does not spend time waiting for an I/O operation to be performed
- ✓ Direct Memory Access
  - blocks of data are moved into memory without involving the processor
  - processor involved at beginning and end only

# Evolution of the I/O Function

## ✓ I/O channel

- I/O module is a separate processor
- Uses computer's main memory

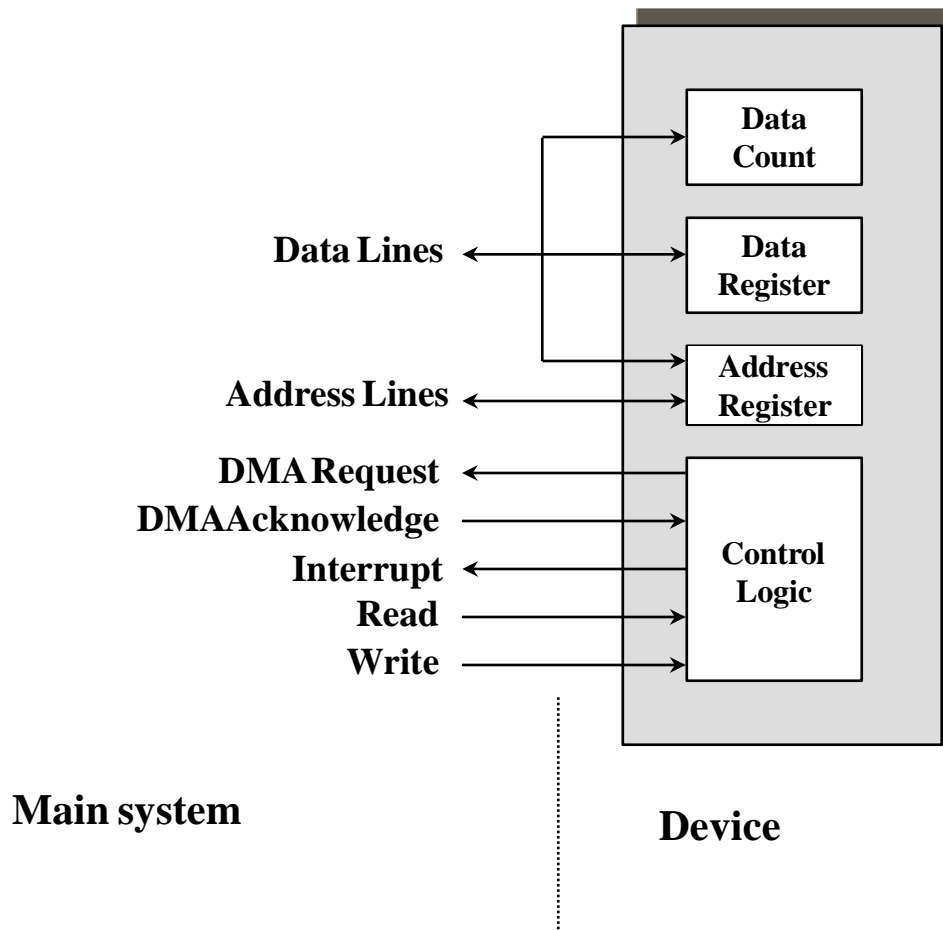
## ✓ I/O processor

- I/O module is a processor with its own local memory
- It's a computer in its own right

# Direct Memory Access

- ✓ Takes control of the system from the CPU to transfer data to and from memory over the system bus
- ✓ Cycle stealing is used to transfer data on the system bus
- ✓ The instruction cycle is suspended so data can be transferred
- ✓ The CPU pauses one bus cycle
- ✓ No interrupts occur
  - does not need to save context

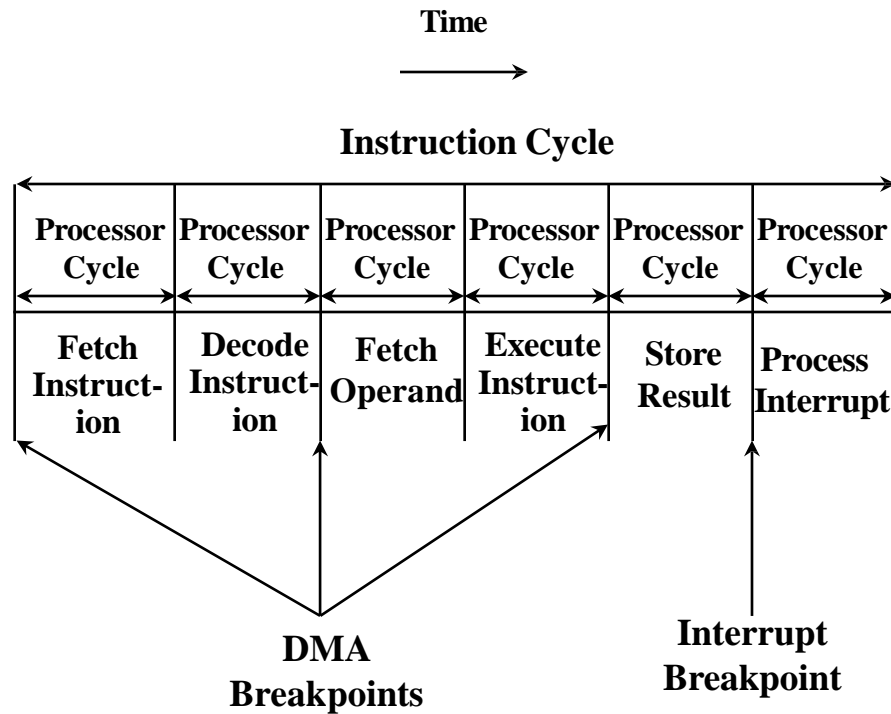
# Typical DMA Block Diagram



# Direct Memory Access


- ✓ Cycle stealing causes the CPU to execute more slowly
- ✓ Number of required busy cycles can be cut by integrating the DMA and I/O functions
- ✓ Try to use path between DMA module and I/O module that does not include the system bus

# DMA and Interrupt Breakpoints




**Breakpoints where CPU can be suspended to let the DMA module use the buss**

# Single-bus, Detached DMA

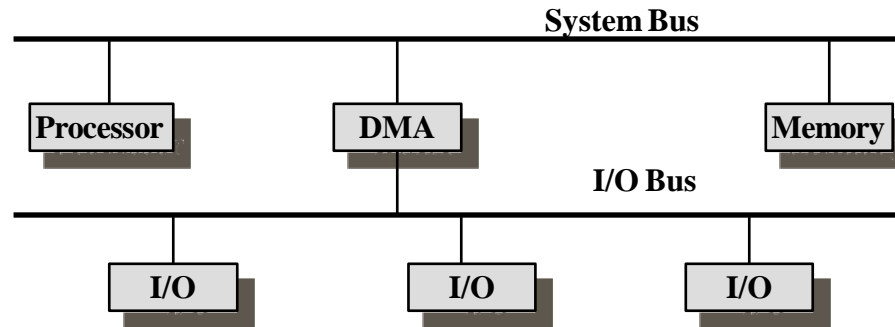




# Single-bus, Integrated DMA-I/O



# I/O Bus



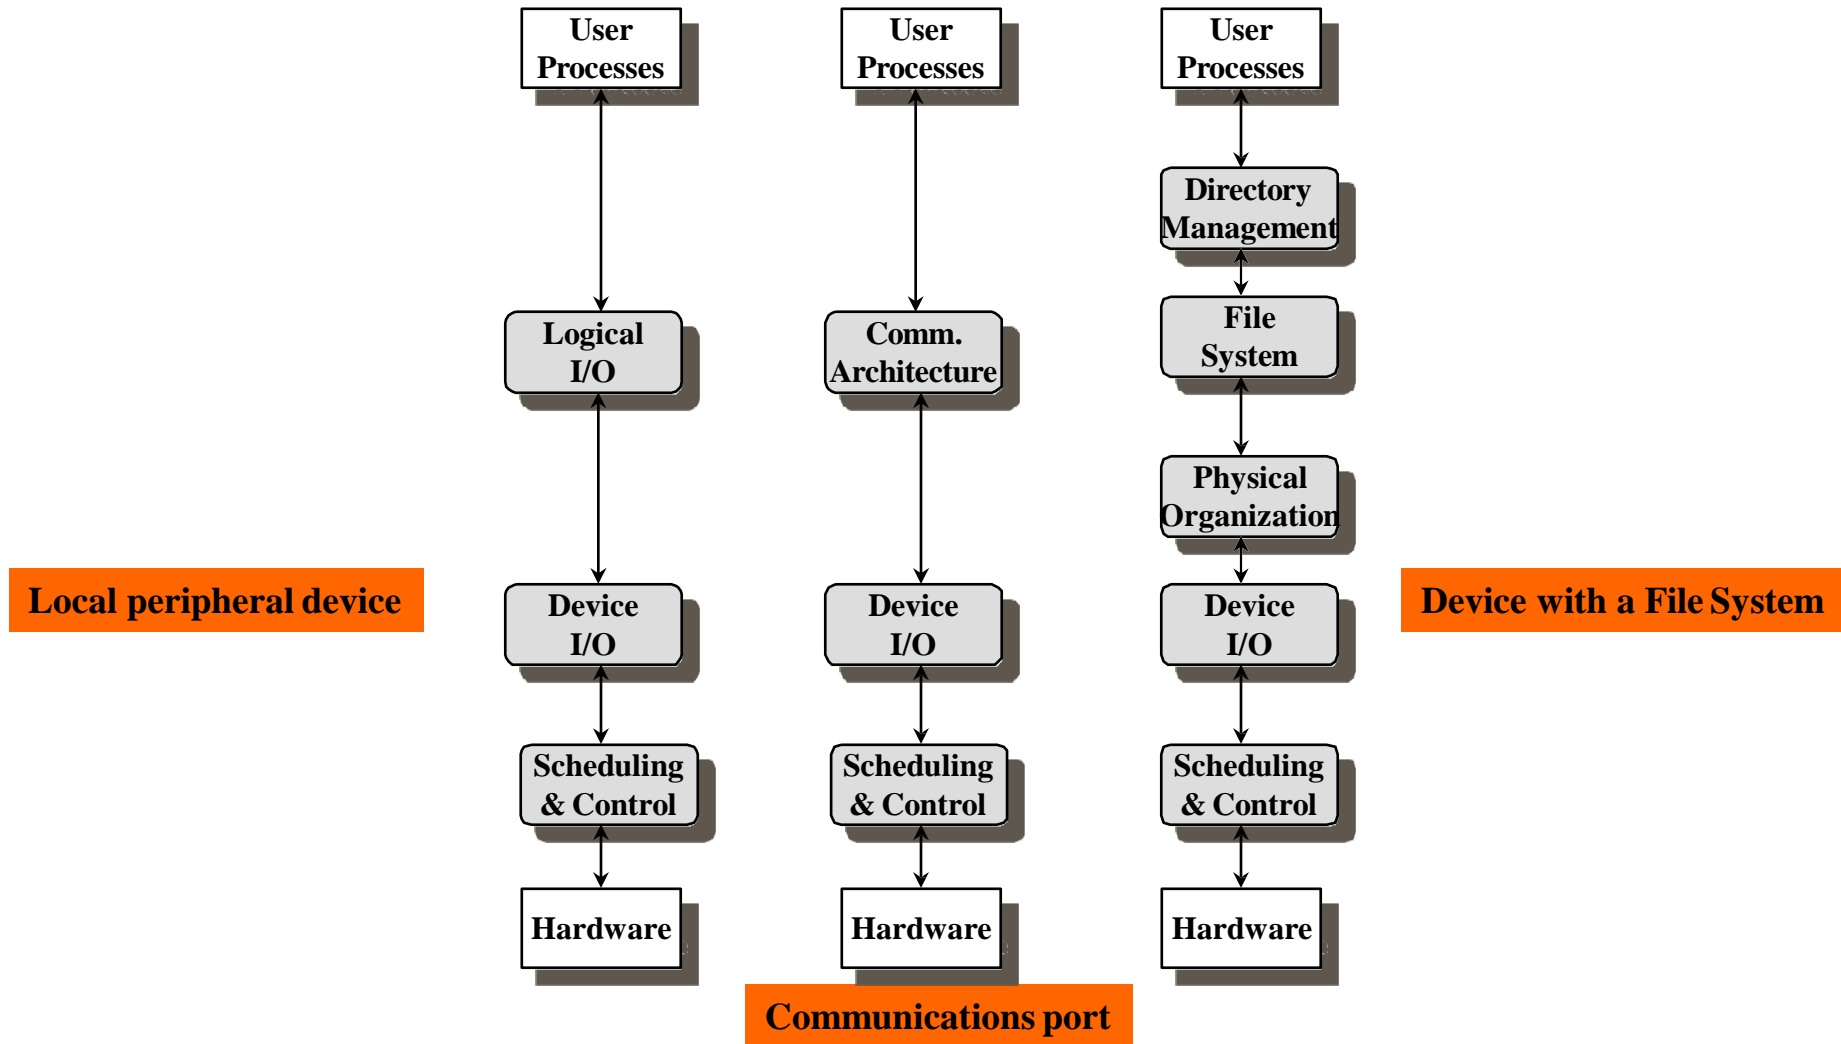
# Operating System Design Objectives

- ✓ I/O is extremely slow compared to main memory
- ✓ Use of multiprogramming allows that some processes will be waiting on I/O while another process executes
- ✓ I/O cannot keep up with processor speed
- ✓ Swapping is used to bring in additional Ready processes, which is an I/O operation
- ✓ Efficiency of I/O is an important issue, since this is a bottleneck

# Operating System Design Objectives

- ✓ Desirable to handle all I/O devices in a uniform manner
- ✓ Hide most of the details of device I/O in lower-level routines so that processes and upper levels see devices in general terms such as Read, Write, Open, and Close
- ✓ Generality is an important issue

# A Model of I/O Organization



# I/O Buffering

- ✓ Reasons for buffering: to find a solution to these problems:
  - Processes must wait for I/O to complete before proceeding
  - Certain pages must remain in main memory during I/O – interferes with page replacement

# I/O Buffering


## ✓ Block-oriented

- information is stored in fixed sized blocks
- transfers are made a block at a time
- used for disks and tapes

## ✓ Stream-oriented

- transfer information as a stream of bytes
- used for terminals, printers, communication ports, mouse, and most other devices that are not secondary storage


# No Buffering





# Single Buffer

- ✓ Operating system assigns a buffer in main memory for an I/O request
- ✓ Block-oriented
  - input transfers are made to buffer
  - block moved to user space when needed
  - another block is moved into the buffer
    - *read ahead*



# Single Buffer

## ✓ *Block-oriented I/O:*

- user process can work on one block of data while next block is being read in
- process waiting for I/O can be swapped out, since input is taking place in system memory, not user memory
- operating system keeps track of assignment of system buffers to user processes
- output is accomplished by the user process writing a block to the buffer and later actually written out


# Single Buffer

## ✓ *Stream-oriented:*

- used one line at a time
- user input from a terminal is one line at a time with carriage return signaling the end of the line
- output to the terminal is one line at a time

# Double Buffer

- ✓ Use two system buffers instead of one
- ✓ A process can transfer data to or from one buffer while the operating system empties or fills the other buffer



# Circular Buffer

- ✓ More than two buffers are used
- ✓ Each individual buffer is one unit in a circular buffer
- ✓ Used when I/O operation must keep up with process

